

Egor Zverev

✉ egor.zverev@ist.ac.at | 🏠 egozverev.github.io/ | 🌐 egozverev | 💻 egor-zverev-ai | 🎓 Egor Zverev

Education

ISTA (Institute of Science and Technology Austria)

Vienna, Austria

PhD in Machine Learning

Sep. 2022 – Present

- Research focus: LLM Security, Instruction-Data Separation, Prompt Injection, AI Agent Security.
- Supervisor: Christoph Lampert (ISTA).
- Joint supervision with Florian Tramèr (ETH Zürich) as part of the ELLIS PhD Program (Sep. 2025 – Present).

MIPT (Moscow Institute of Physics and Technology)

Moscow, Russia

B.S. in Applied Mathematics and Computer Science, Yandex Department of Data Analysis

Aug. 2018 – Aug. 2022

- Relevant courses: Machine Learning, NLP, Statistics, Bayesian Methods, Optimization, C++, Python, Algorithms.
- Excellence Scholarship (for **top-7%** 4th-year students) and Abramov Scholarship (for **top-10%** junior students).
- GPA: 4.8/5.0 (graduated **with honors**).

Publications

Untrusted Content Masking for Web Agents with Security Guarantees

Kristina Nikolić*, Egor Zverev*, Javier Rando, Matthew Jagielski, Edoardo Debenedetti, Florian Tramèr

Advances in Neural Information Processing Systems (NeurIPS), 2026, *Equal contribution URL: <https://arxiv.org/abs/2607.05277>

ASIDE: Architectural Separation of Instructions and Data in Language Models

Egor Zverev, Evgenii Kortukov, Alexander Panfilov ... Sebastian Lapuschkin, Wojciech Samek, Christoph H. Lampert

International Conference on Learning Representations (ICLR), 2026, URL: <https://arxiv.org/abs/2503.10566>

Can LLMs Separate Instructions From Data? And What Do We Even Mean By That?

Egor Zverev, Sahar Abdelnabi, Soroush Tabesh, Mario Fritz, Christoph H. Lampert

International Conference on Learning Representations (ICLR), 2025, URL: <https://arxiv.org/abs/2403.06833>

LLMail-Inject: A Dataset from a Realistic Adaptive Prompt Injection Challenge

Sahar Abdelnabi, Aideen Fay, Ahmed Salem, Egor Zverev, Kai-Chieh Liao ... Andrew Pavard, Giovanni Cherubin

Findings of the Association for Computational Linguistics (EMNLP), 2026, URL: <https://arxiv.org/abs/2506.09956>

Work Experience

Microsoft Research Cambridge

Cambridge, UK

Research Intern (Privacy & Security)

Jun. 2026 – Sep. 2026

- Developed a method for enforcing Information-Flow Control in Browser Agents (paper under submission).

ETH Robotics Club

Zürich, Switzerland

Research Advisor

Mar. 2026 – Present

- Supervising several master's students (Andrei Baroian, Noé Macé, and Gustave Charles Saigne) on a project exploring AI security vulnerabilities in state-of-the-art robotics models.

ETH Zürich

Zürich, Switzerland

Academic Guest

Sep. 2025 – Dec. 2025

- Visited Florian Tramèr's lab as part of the ELLIS PhD Program to work on Untrusted Content Masking.

MTS Digital

Moscow, Russia

Junior Data Scientist at Experimental Group (A/B + R&D)

Jan. 2022 – Jun. 2022

- Sped up statistical experiments for movie recommendation systems from 3 to 2 weeks by using machine learning techniques.
- Contributed to an A/B testing library used by 70+ data scientists.

Google Summer of Code (MapAction)

Remote (UK)

Student Developer

May 2021 – Aug. 2021

- Awarded a **Google stipend** to develop an open-source project on humanitarian data acquisition.
- Extended an Airflow ETL pipeline to cover 10 data products (up from 3); redesigned the code structure and reduced the codebase by 30%.

Leadership & Service

ml4good AI Governance Bootcamp

Paris, France

Technical Teacher

Sep. 2026

- Developed from scratch a series of 6 lectures and exercises on foundations of ML and AI Safety for non-technical audiences.
- Taught the course to AI Governance students at the Sep. 2026 ml4good bootcamp.

Foundations of Language Models Security Workshop @ NeurIPS 2026

Lead Organizer

Paris, France

Jul. 2026 – Present

- Leading organization of the second iteration of the FLMSEC workshop.

Foundations of Language Models Security Workshop @ EurIPS 2025

Lead Organizer

Copenhagen, Denmark

Aug. 2025 – Dec. 2025

- Proposed a new workshop format prioritizing community building; gathered a team of 5 co-organizers and coordinated their work.
- Created the [workshop website](#), managed OpenReview submissions, and led onsite organization.

LLM Safety and Security Workshop @ ELLIS UnConference 2025

Lead Organizer

Copenhagen, Denmark

Oct. 2025 – Dec. 2025

- Gathered an organizing team, managed the [website and submissions](#), and led onsite organization.

ISTA AI Safety Discussion Group

Organizer

Vienna, Austria

Feb. 2023 – May 2023

- Organized a weekly AI Safety paper reading group.

MIPT.Stats

Statistics Teaching Assistant

Moscow, Russia

Aug. 2021 – Feb. 2022

- Conducted weekly one-on-one meetings with 12 3rd-year statistics students and checked their practical homework.

Invited Talks & Panels

Secure-by-design AI Agents: Threats, Guarantees, and Global Security	Princeton School on Science and Global Security	Jul. 2026
Untrusted Content Masking for Web Agents with Security Guarantees	Imperial College London	Jul. 2026
Science Tour 2026: a series of talks on “Untrusted Content Masking” and ASIDE papers	ELLIS Institute Tübingen, Saarland University, University of Luxembourg, University of Liechtenstein, CISPA	Mar. 2026
Co-panelist, “Real-time Safeguards for Human–AI Interaction” track (recording)	IASEAI’26, UNESCO, Paris	Feb. 2026
Instruction-Data Separation	IBM Research, San Jose	Sep. 2025
ASIDE: Architectural Separation of Instructions and Data in LLMs	BuildTrust @ ICLR 2025, Singapore	Apr. 2025
ASIDE: Architectural Separation of Instructions and Data in LLMs	ETH Zürich	Mar. 2025

Professional Development

Summer Schools

Princeton School on Science and Global Security 2026 (talk), CISPA-ELLIS Trustworthy AI 2025 (poster), EEML 2024 (poster), Cooperative AI 2023

AI Safety Courses

AISES 2024, Intro to ML Safety 2023, AI Safety Fundamentals 2022