

Egor Zverev

✉ egor.zverev@ist.ac.at | 🏠 egozverev.github.io/ | 🐙 egozverev | 💻 egor-zverev-ai | 🎓 Egor Zverev

Education

ISTA (Institute of Science and Technology Austria)

PhD in Machine Learning

Vienna, Austria

Sep. 2022 – Present

- Research focus: LLM Security, Instruction-Data Separation, Prompt Injection, AI Agent Security.
- Supervisor: [Christoph Lampert](#) (ISTA).
- Joint supervision with [Florian Tramèr](#) (ETH Zürich) as part of the ELLIS PhD Program (Sep. 2025 – Present).

MIPT (Moscow Institute of Physics and Technology)

B.S. in Applied Mathematics and Computer Science, Yandex Department of Data Analysis

Moscow, Russia

Aug. 2018 – Aug. 2022

- Relevant courses: Machine Learning, NLP, Statistics, Bayesian Methods, Optimization, C++, Python, Algorithms.
- Excellence Scholarship (for **top-7%** 4th-year students) and Abramov Scholarship (for **top-10%** junior students).
- GPA: 4.8/5.0 (graduated **with honors**).

Publications

Untrusted Content Masking for Web Agents with Security Guarantees

Kristina Nikolić*, Egor Zverev*, Javier Rando, Matthew Jagielski, Edoardo Debenedetti, Florian Tramèr

ICML Workshop on Agents in the Wild (AIWILD), Spotlight, 2026, *Equal contribution URL: <https://openreview.net/forum?id=u1YuI0K3eP>

ASIDE: Architectural Separation of Instructions and Data in Language Models

Egor Zverev, Evgenii Kortukov, Alexander Panfilov ... Sebastian Lapuschkin, Wojciech Samek, Christoph H. Lampert

International Conference on Learning Representations (ICLR), 2026, URL: <https://arxiv.org/abs/2503.10566>

Can LLMs Separate Instructions From Data? And What Do We Even Mean By That?

Egor Zverev, Sahar Abdelnabi, Soroush Tabesh, Mario Fritz, Christoph H. Lampert

International Conference on Learning Representations (ICLR), 2025, URL: <https://arxiv.org/abs/2403.06833>

LLMail-Inject: A Dataset from a Realistic Adaptive Prompt Injection Challenge

Sahar Abdelnabi, Aideen Fay, Ahmed Salem, Egor Zverev, Kai-Chieh Liao ... Andrew Paverd, Giovanni Cherubin

EurIPS, 2025, URL: <https://arxiv.org/abs/2506.09956>

Work Experience

Microsoft Research Cambridge

Research Intern (Privacy & Security)

Cambridge, UK

Jun. 2026 – Present

- Working with Santiago Zanella-Béguelin on secure-by-design web agents.

ETH Robotics Club

Research Advisor

Zürich, Switzerland

Mar. 2026 – Present

- Supervising several master's students (Andrei Baroian, Noé Macé, and Gustave Charles Saigne) on a project exploring AI security vulnerabilities in state-of-the-art robotics models.

ETH Zürich

Academic Guest

Zürich, Switzerland

Sep. 2025 – Dec. 2025

- Visited Florian Tramèr's lab as part of the ELLIS PhD Program to work on Untrusted Content Masking.

MTS Digital

Junior Data Scientist at Experimental Group (A/B + R&D)

Moscow, Russia

Jan. 2022 – Jun. 2022

- Sped up statistical experiments for movie recommendation systems from 3 to 2 weeks by using machine learning techniques.
- Contributed to an A/B testing library used by 70+ data scientists.

Google Summer of Code (MapAction)

Student Developer

Remote (UK)

May 2021 – Aug. 2021

- Awarded a **Google stipend** to develop an open-source project on humanitarian data acquisition.
- Extended an Airflow ETL pipeline to cover 10 data products (up from 3); redesigned the code structure and reduced the codebase by 30%.

Leadership & Service

Foundations of Language Models Security Workshop @ NeurIPS 2026

Lead Organizer

Paris, France

Jul. 2026 – Present

- Leading organization of the second iteration of the FLMSEC workshop.

Foundations of Language Models Security Workshop @ EurIPS 2025 Lead Organizer - Proposed a new workshop format prioritizing community building; gathered a team of 5 co-organizers and coordinated their work. - Created workshop website, managed OpenReview submissions, and led onsite organization.	Copenhagen, Denmark Aug. 2025 – Dec. 2025
LLM Safety and Security Workshop @ ELLIS UnConference 2025 Lead Organizer Gathered an organizing team, managed the website and submissions, and led onsite organization.	Copenhagen, Denmark Oct. 2025 – Dec. 2025
ISTA AI Safety Discussion Group Organizer Organized a weekly AI Safety paper reading group.	Vienna, Austria Feb. 2023 – May 2023
MIPT.Stats Statistics Teaching Assistant Conducted weekly one-on-one meetings with 12 3rd-year statistics students and checked their practical homework.	Moscow, Russia Aug. 2021 – Feb. 2022

Invited Talks & Panels

Secure by Design AI Agents: Threats, Guarantees, and Global Security	Princeton School on Science and Global Security	Jul. 2026
Untrusted Content Masking for Web Agents with Security Guarantees	Imperial College London	Jul. 2026
Science Tour 2026: a series of talks on “Untrusted Content Masking” and ASIDE papers	ELLIS Institute Tübingen, Saarland University, University of Luxembourg, University of Liechtenstein, CISPA	Mar. 2026
Co-panelist, “Real-time Safeguards for Human–AI Interaction” track (recording)	IASEAI’26, UNESCO, Paris	Feb. 2026
Instruction-Data Separation	IBM Research, San Jose	Sep. 2025
ASIDE: Architectural Separation of Instructions and Data in LLMs	BuildTrust @ ICLR 2025, Singapore	Apr. 2025
ASIDE: Architectural Separation of Instructions and Data in LLMs	ETH Zürich	Mar. 2025

Professional Development

Summer Schools	Princeton School on Science and Global Security 2026 (talk), CISPA-ELLIS Trustworthy AI 2025 (poster), EEML 2024 (poster), Cooperative AI 2023
AI Safety Courses	AISES 2024, Intro to ML Safety 2023, AI Safety Fundamentals 2022